

Advanced Database Systems (CS60113)

Assignment 5. In-Memory Databases

(Due Date: November 3 2020)

NOTE: This is an individual assignment.

Create two versions of the following four tables in MySQL: One for disk resident database and one for In-memory database. You may name those with a prefix `dsk` and `inm` for convenience. You will need to change the table names in the queries (given below) accordingly. For creating a table as In-memory, use the following clause in the create table script: `ENGINE=MEMORY`. Note, any data inserted in In-memory tables will go once you restart your machine

`storedim` (`scode` integer, `sname` varchar(30), `city` varchar(30), `state` varchar(20), `pincode` integer)
`itemdim` (`icode` integer, `iname` varchar(30), `category` char(5), `subcategory` char(3), `price` integer)
`dateofpurchasedim` (`dcode` integer, `purchasedate` date, `dayoftheweek` char(3), `month` char(3), `quarter` char(2), `year` integer)
`purchasefact` (`scode`, `icode`, `dcode`, `qty` integer). `scode`, `icode`, `dcode` are foreign keys from the previous three tables

Descriptions of columns and possible values for the relevant ones are given below.

`scode`, `icode`, `dcode`, auto-generated unique serial numbers. You may also generate on your own.

`sname`, `iname`: Random strings

`city`: Mumbai, Chennai, Kolkata, Pune, Bengaluru, Guwahati, Kochi

`state`: Kerala, Maharashtra, West Bengal, Karnataka, Tamil Nadu, Assam (Ensure right city is included under right state)

`pincode`: 123456, 234567, 345678, 456789, 567890, 678901, 789012, 890123, 901234. No constraint with respect to city and state.

`category`: GROCR, CONSM, FASHN, MEDCN, VGTBL

`subcategory`: PRS, OFF, MIL. No relationship with category

`price`: Integer

`purchasedate`: Date type column.

`dayoftheweek`: MON, TUE, WED, THR, FRI, SAT, SUN

`month`: JAN, FEB, ...DEC

`quarter`: Q1, Q2, Q3, Q4 (Ensure correct matching of months to quarter, i.e., APR, MAY, JUN under Q1, JUL, AUG, SEP under Q2, etc.

`year`: 2015, 2016, 2017, 2018, 2019, 2020

Write programs in any language to generate appropriate "Insert into" statements for the 4 tables. Note, exactly the same rows should be entered in the two sets of tables – Disk resident and In-memory.

Generate 1000 rows each for `storedim`, `itemdim` and `dateofpurchasedim`.

Generate rows for `purchasefact` (100, 1,000, 100,000, 1,000,000, 10,000,000 number of rows in steps).

Save the two sets (one for each type of DB) of 4 "insert into" scripts (3 for the four dimension tables + 1 for entering 10,000,000 rows in the fact table). You will need to submit these for the assignment. Also, you will be able to use these scripts if you restart your machine and the In-memory data will be gone.

Execute the following queries:

Q1

Select s.state, s.city, s.pincode, sum(p.qty) from storedim s, purchasefact p where s.scode = p.scode group by s.state, s.city, s.pincode

Q2

Select s.sname, i.iname, d. purchasedate, p.qty from storedim s, itemdim i, dateofpurchasedim d, purchasefact p where s.scode = p.scode and i.icode=p.icode and d.dcode = p.dcode

Q3

Select s.city, i.category, sum(p.qty*i.price) from storedim s, itemdim i, purchasefact p where s.scode = p.scode and i.icode=p.icode group by s.city, i.category

Q4

Select s.city, i.subcategory, d.month, sum(p.qty*i. price) from storedim s, itemdim i, dateofpurchasedim d, purchasefact p where s.scode = p.scode and i.icode=p.icode and d.dcode = p.dcode group by s.city, i.subcategory, d.month

Q5

Select s.state, s.city, s.pincode, i.category, i.subcategory, d.year, d.quarter, d.month,d.dayoftheweek, sum(p.qty*i. price) from storedim s, itemdim i, dateofpurchasedim d, purchasefact p where s.scode = p.scode and i.icode=p.icode and d.dcode = p.dcode group by s.state, s.city, s.pincode, i.category, i.subcategory, d.year, d.quarter, d.month, d.dayoftheweek

Note down the time (in seconds) for executing the queries. Report your results in the following tabular form.

	100 rows	1,000 rows	100,000 rows	1,000,000 rows	10,000,000 rows
Q1	T1/T2	T1/T2	T1/T2	T1/T2	T1/T2
Q2	T1/T2	T1/T2	T1/T2	T1/T2	T1/T2
Q3	T1/T2	T1/T2	T1/T2	T1/T2	T1/T2
Q4	T1/T2	T1/T2	T1/T2	T1/T2	T1/T2
Q5	T1/T2	T1/T2	T1/T2	T1/T2	T1/T2

T1: Time for Disk resident database (in seconds)

T2: Time for In-memory database (in seconds)

Submit (i) a zip file containing the eight create table scripts (4 for Disk resident + 4 for In-memory) and eight insert into scripts (4 for Disk resident + 4 for In-memory) and (ii) a pdf file containing the results in the above tabular format. Note that, you do not need to submit the programs for generating the insert into scripts. Note, I should be able to simply execute these queries to recreate your environment and re-run the experiments if required.